Prompt 1 · Orientation: Six Industries, Cold Read

Six case cards, six money questions, and one honest note on where the cold read was thin.

A · Café Operations *(cafe_sales.csv)*

A multi-site café's point-of-sale log. One row = one item sold in one transaction (product, daypart, payment method). Span: Jun 2025 - May 2026 (~6,159 rows). **Contested column: total_eur**, since revenue looks strong but says nothing about margin after card-processing fees.
Stakeholder: Store Ops Manager. Money question: which dayparts/products actually drive revenue once payment-method mix is accounted for, and is weekend/summer lift big enough to justify extra staffing?

B · SaaS Support *(support_tickets.csv)*

A SaaS company's support desk log. One row = one ticket, from open to resolution, with category, priority, sentiment and CSAT. Span: Dec 2025 - May 2026 (600 rows). **Contested column: customer_sentiment**, a subjective label that drives churn scoring but can disagree with the ticket's actual tone.
Stakeholder: Head of Customer Support. Money question: which categories/channels combine slow resolution with negative sentiment - i.e. where is churn risk concentrated?

C · E-commerce Reviews *(product_reviews.csv)*

Reviews for a coffee/kitchenware retailer's products. One row = one customer review (rating, text, verified flag). Span: Jun 2025 - May 2026 (500 rows). **Contested column: verified_purchase**, since it decides whether a complaint counts as real signal or noise.
Stakeholder: Head of Product. Money question: which complaint themes (taste, delivery, price, packaging) are costly enough to justify a fix, and should unverified reviews count the same as verified ones?

D · Banking & Credit *(loan_applications.csv)*

A lender's application-to-decision file. One row = one loan application with applicant profile and underwriting decision. Span: Sep 2025 - May 2026 (800 rows, 50.4% approved). **Contested column: region**, an innocuous-looking field that turns out to drive outcomes inside the discretionary score band.
Stakeholder: Chief Risk / Compliance Officer. Money question: is the ~50% approval rate explained by risk alone, or does region/credit-band interaction expose the lender to a disparate-impact claim?

E · Finance (AP) *(invoices.csv)*

A company's accounts-payable ledger. One row = one vendor invoice, issue-to-payment. Span: Jul 2025 - Jun 2026 due dates (603 rows; 149 overdue, 1 exact duplicate found). **Contested column: status**, since overdue vs. pending vs. paid determines how much cash looks at risk.
Stakeholder: CFO / AP Manager. Money question: how much cash is genuinely at risk in 'overdue', versus invoices that are simply missing a payment_date entry?

F · Marketing *(marketing_campaigns.csv)*

Daily channel-level paid/organic/email performance. One row = one channel-campaign-day (spend, clicks, conversions, revenue, ROAS). Span: Jan-May 2026 only (750 rows). **Contested column: roas**, since it looks like a clean return metric but depends entirely on the attribution model behind it.
Stakeholder: Marketing Director. Money question: where should the next budget euro go, and is the channel with implausibly high ROAS a real star or an attribution artifact?

Where the AI's answer was thin

Before trusting the six case cards, I checked what date range each file actually covers. The AI's cold read treated all six files as if they belonged to roughly the same period, since they came from the same exercise batch. They don't. Four files (cafe_sales, invoices, loan_applications, product_reviews) span close to a full year, June/July 2025 through May 2026. The other two are much narrower: support_tickets only starts in December 2025, and marketing_campaigns starts exactly on January 1, 2026, with nothing before it.

That last one is the trap the AI would fall into if I let it. It would be easy for it to look at marketing_campaigns and support_tickets side by side and claim something like "marketing spend and café traffic both climbed in Q1 2026" - a sentence that sounds like a genuine insight. But for marketing_campaigns, Q1 2026 isn't a period where spend rose; it's the first data point that exists at all. There's nothing before it to compare against, so any statement about a trend in that window would really just be an artifact of missing history, not a real pattern.

What I checked instead: the actual min/max date for every file, before accepting any cross-file comparison the AI makes later in the exercise. Mismatched time windows are a hard constraint on what can be compared - not a footnote to add after the fact.